

Case Study: NetSuite Data Pipeline

Background: The CFO and Accounting Team needed transactional data from Stripe loaded into NetSuite.

Challenges: NetSuite has a latency of .75s/Journal Entry. After close, data can no longer be changed. Stripe can change transactional data anytime within a 45 day window.

Solution: Build a data pipeline using AWS Lambda, Batch, S3, and Snowflake for staging into NetSuite and Tableau.

Coding Stack: Python, SQL, and Spark.

Tradeoffs: Batch vs. streaming – Stripe transactional data changes intermittently, so batch was more appropriate. Boomi vs. in-house Python SDK – the accounting team needed an ETL tool that they could monitor, so we built both solutions.

Security and dev ops: ACL for each stage of the process with one VPC for the AWS pipeline.

End result: Successful first close.

Architectural Diagram

